



databricks

(https://databricks.com) Python



Import notebook

1

```
# List only datasets larger than 1 GB
def list_large_datasets(path="/databricks-datasets", max_depth=2, current_depth=0):
    if current_depth >= max_depth:
        return
    try:
        files = dbutils.fs.ls(path)
        for f in files:
            if f.isDir():
                # Compute folder size
                try:
                    sub_files = dbutils.fs.ls(f.path)
                    total_size = sum(sf.size for sf in sub_files)
                    size_gb = total_size / (1024 ** 3)
                    if size_gb > 1:
                        print(f"{f.path} --> {size_gb:.2f} GB")
                    # Recurse deeper
                    list_large_datasets(f.path, max_depth, current_depth + 1)
                except Exception:
                    pass
    except Exception as e:
        print(f"Error accessing {path}: {e}")

# Run it
list_large_datasets()
```

```
dbfs:/databricks-datasets/COVID/CORD-19/ --> 7.00 GB
dbfs:/databricks-datasets/airlines/ --> 120.06 GB
dbfs:/databricks-datasets/asa/airlines/ --> 11.20 GB
dbfs:/databricks-datasets/genomics/grch37/ --> 1.65 GB
dbfs:/databricks-datasets/genomics/grch37_merged_vep_96/ --> 13.28 GB
dbfs:/databricks-datasets/genomics/grch37_refseq_vep_96/ --> 10.94 GB
dbfs:/databricks-datasets/genomics/grch37_star/ --> 27.85 GB
dbfs:/databricks-datasets/genomics/grch37_vep/ --> 13.28 GB
dbfs:/databricks-datasets/genomics/grch37_vep_96/ --> 11.50 GB
dbfs:/databricks-datasets/genomics/grch38/ --> 1.57 GB
dbfs:/databricks-datasets/genomics/grch38_merged_vep_96/ --> 13.91 GB
```

```
dbfs:/databricks-datasets/genomics/grch38_refseq_vep_96/ --> 11.17 GB  
dbfs:/databricks-datasets/genomics/grch38_star/ --> 33.43 GB  
dbfs:/databricks-datasets/genomics/grch38_vep/ --> 14.67 GB  
dbfs:/databricks-datasets/genomics/grch38_vep_96/ --> 11.88 GB  
dbfs:/databricks-datasets/learning-spark-v2/sf-fire/ --> 1.07 GB  
dbfs:/databricks-datasets/med-images/camelyon16/ --> 109.81 GB  
dbfs:/databricks-datasets/sai-summit-2019-sf/ --> 1.06 GB  
dbfs:/databricks-datasets/timeseries/Fires/ --> 1.76 GB  
dbfs:/databricks-datasets/wiki/ --> 4.41 GB
```

```
# Simple Databricks cell: get dataset size and shape  
  
path = 'dbfs:/databricks-datasets/timeseries/Fires/'  
  
df = spark.read.option("header", True).csv(path)  
rows = df.count()  
cols = len(df.columns)  
  
# get total size in GB  
size_bytes = sum(f.size for f in dbutils.fs.ls(path))  
size_gb = size_bytes / (1024 ** 3)  
  
print(f"Rows: {rows}")  
print(f"Columns: {cols}")  
print(f"Size: {size_gb:.3f} GB")
```

```
▶ 📄 df: pyspark.sql.connect.dataframe.DataFrame = [Call Number: string, Unit ID: string ...  
32 more fields]
```

```
Rows: 5120231  
Columns: 34  
Size: 1.763 GB
```

```
df.printSchema()
```

```
|-- Available DtTm: string (nullable = true)  
|-- Address: string (nullable = true)  
|-- City: string (nullable = true)
```

```
|-- Zipcode of Incident: string (nullable = true)
|-- Battalion: string (nullable = true)
|-- Station Area: string (nullable = true)
|-- Box: string (nullable = true)
|-- Original Priority: string (nullable = true)
|-- Priority: string (nullable = true)
|-- Final Priority: string (nullable = true)
|-- ALS Unit: string (nullable = true)
|-- Call Type Group: string (nullable = true)
|-- Number of Alarms: string (nullable = true)
|-- Unit Type: string (nullable = true)
|-- Unit sequence in call dispatch: string (nullable = true)
|-- Fire Prevention District: string (nullable = true)
|-- Supervisor District: string (nullable = true)
|-- Neighborhoods - Analysis Boundaries: string (nullable = true)
|-- Location: string (nullable = true)
|-- RowID: string (nullable = true)
```

```
# Select important columns from Fire Calls dataset
df_spark_select = df.select(
    "Call Type",
    "City",
    "Response DtTm",
    "On Scene DtTm"
)

print("PySpark:")
df_spark_select.show(5)
```

▶ 📈 df_spark_select: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 2 more fields]

PySpark:

Call Type	City	Response DtTm	On Scene DtTm
Alarms	San Francisco	10/18/2019 12:08:...	10/18/2019 12:11:...
Alarms	San Francisco	NULL	NULL
Alarms	San Francisco	NULL	NULL
Alarms	San Francisco	10/18/2019 12:09:...	10/18/2019 12:09:...

```
|Structure Fire|San Francisco|10/18/2019 12:13:...|10/18/2019 12:16:...|  
+-----+-----+-----+-----+  
only showing top 5 rows
```

```
from pyspark.sql.functions import col  
  
# Apply multiple filters on Fire Calls dataset  
df_filtered = (  
    df_spark_select  
    # Filter 1: Keep only rows where Call Type is "Structure Fire"  
    .filter(col("Call Type") == "Structure Fire")  
    # Filter 2: Keep only rows where City is "San Francisco"  
    .filter(col("City") == "San Francisco")  
)  
  
print("✅ PySpark – Filtered Fire Calls (Structure Fires in San  
Francisco):")  
df_filtered.show(5, truncate=False)
```

```
▶ 📄 df_filtered: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string  
... 2 more fields]  
✅ PySpark – Filtered Fire Calls (Structure Fires in San Francisco):  
+-----+-----+-----+-----+  
-+  
|Call Type      |City           |Response DtTm          |On Scene DtTm  
|  
+-----+-----+-----+-----+  
-+  
|Structure Fire|San Francisco|10/18/2019 12:13:52 AM|10/18/2019 12:16:16 A  
M|  
|Structure Fire|San Francisco|10/18/2019 12:14:28 AM|NULL  
|  
|Structure Fire|San Francisco|10/18/2019 02:24:56 AM|10/18/2019 02:28:08 A  
M|  
|Structure Fire|San Francisco|09/17/2017 10:06:29 AM|NULL  
|  
|Structure Fire|San Francisco|10/18/2019 07:08:33 AM|NULL  
|  
+-----+-----+-----+-----+
```

```
-+  
only showing top 5 rows
```

```
# Aggregation on the Fire Calls data  
from pyspark.sql.functions import col, to_timestamp, avg, count,  
min as min_, max as max_, desc, expr  
  
TS_FMT = "M/d/yyyy h:mm:ss a"  
  
# 1) Add response delay (minutes) from the selected columns  
df_with_delay = (  
    df_spark_select  
        .withColumn("received_ts", to_timestamp(col("Response DtTm"),  
TS_FMT))  
        .withColumn("on_scene_ts", to_timestamp(col("On Scene DtTm"),  
TS_FMT))  
        .withColumn("response_delay_min",  
(col("on_scene_ts").cast("long") -  
col("received_ts").cast("long")) / 60.0)  
)  
  
# 2) Aggregation per Call Type, compute count, avg, median, p95,  
min, max delays  
agg_result = (  
    df_with_delay  
        .filter(col("response_delay_min").isNotNull() &  
(col("response_delay_min") > 0))  
        .groupBy("Call Type")  
        .agg(  
            count("*").alias("count_calls"),  
            avg("response_delay_min").alias("avg_delay_min"),  
            expr("percentile_approx(response_delay_min,  
0.5)").alias("median_delay_min"),  
            expr("percentile_approx(response_delay_min,  
0.95)").alias("p95_delay_min"),  
            min_("response_delay_min").alias("min_delay_min"),  
            max_("response_delay_min").alias("max_delay_min"),  
)  
        .orderBy(desc("p95_delay_min"))
```

```

    .limit(10)
)

agg_result.show(truncate=False)

```

▶ 📄 agg_result: pyspark.sql.connect.DataFrame = [Call Type: string, count_calls: long ... 5 more fields]

▶ 📄 df_with_delay: pyspark.sql.connect.DataFrame = [Call Type: string, City: string ... 5 more fields]

```

+
|Administrative |66 |27.54848484848485 |14.03
33333333333333|90.25 |0.05 |93.16666666666667
|
|Mutual Aid / Assist Outside Agency|173 |28.03150289017341 |12.61
6666666666667|77.4 |0.0333333333333333 |265.25
|
|Aircraft Emergency |507 |19.527843523997372 |14.35
|49.81666666666667 |0.05 |119.0666666666666 |
|Watercraft in Distress |457 |9.325054704595184 |5.916
6666666666667 |30.25 |0.1333333333333333 |116.51666666666667
|
|Train / Rail Incident |885 |7.62864406779661 |4.316
666666666666 |26.83333333333332 |0.0166666666666666 |116.1
|
|Water Rescue |12277 |8.685153810648638 |6.0
|24.25 |0.0166666666666666 |226.96666666666667 |
|Suspicious Package |220 |7.390757575757576 |4.8
|22.75 |0.0166666666666666 |43.15 |
|High Angle Rescue |788 |9.081641285956005 |5.966
666666666667 |20.83333333333332 |0.0333333333333333 |388.36666666666667

```

```

from pyspark.sql.functions import col, avg, count, desc,
to_timestamp

```

```

# GroupBy with aggregation on Fire Calls dataset
TS_FMT = "M/d/yyyy h:mm:ss a"

```

```

# Add response delay in minutes
df_with_delay = (

```

```

        df_spark_select
        .withColumn("received_ts", to_timestamp(col("Response DtTm"),
TS_FMT))
        .withColumn("on_scene_ts", to_timestamp(col("On Scene DtTm"),
TS_FMT))
        .withColumn("response_delay_min",
(col("on_scene_ts").cast("long") -
col("received_ts").cast("long")) / 60.0)
    )

# Group by Call Type and calculate average delay and count of
calls
df_grouped = (
    df_with_delay
    .filter(col("response_delay_min").isNotNull() &
(col("response_delay_min") > 0))
    .groupBy("Call Type")
    .agg(
        avg("response_delay_min").alias("avg_response_delay_min"),
        count("*").alias("count_calls")
    )
    .orderBy(desc("avg_response_delay_min"))
)

print("✅ PySpark – GroupBy with Aggregation:")
df_grouped.show(10, truncate=False)

```

- ▶ 📊 df_grouped: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_response_delay_min: double ... 1 more field]
- ▶ 📊 df_with_delay: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 5 more fields]

✅ PySpark – GroupBy with Aggregation:

Call Type	avg_response_delay_min	count_calls
Mutual Aid / Assist Outside Agency	28.03150289017341	173
Administrative	27.54848484848485	66
Aircraft Emergency	19.527843523997372	507
Watercraft in Distress	9.325054704595184	457
High Angle Rescue	9.081641285956005	788
Water Rescue	8.685153810648638	12277

Train / Rail Incident	7.62864406779661	885	
Suspicious Package	7.3907575757576	220	
Marine Fire	6.943526785714285	224	
Confined Space / Structure Collapse	6.7111986001749795	381	
+-----+-----+-----+			
only showing top 10 rows			

```
from pyspark.sql.functions import col, to_timestamp, round, upper,
concat_ws

# Column transformations using withColumn on Fire Calls dataset
TS_FMT = "M/d/yyyy h:mm:ss a"

df_transformed =
    df_spark_select
    # Convert timestamps
    .withColumn("received_ts", to_timestamp(col("Response DtTm"),
TS_FMT))
    .withColumn("on_scene_ts", to_timestamp(col("On Scene DtTm"),
TS_FMT))
    # Compute new column: response delay in minutes
    .withColumn("response_delay_min",
round((col("on_scene_ts").cast("long") -
col("received_ts").cast("long")) / 60.0, 2))
    # Uppercase city name
    .withColumn("City_Upper", upper(col("City")))
    # Combine city and call type into one string
    .withColumn("Call_Info", concat_ws(" - ", col("City"),
col("Call Type")))
)

print("✅ PySpark – Column Transformations using withColumn:")
df_transformed.select("City", "Call Type", "response_delay_min",
"City_Upper", "Call_Info").show(5, truncate=False)
```

▶ _sqlldf: pyspark.sql.connect.DataFrame = [day_of_week: string, total_calls: long]

Table

ⓘ This result is stored as `_sqldf` and can be used in other Python and SQL cells.

```
%sql
-- Assumes a temp view named `fire_calls` exists with columns:
-- `Call Type`, City, `Response DtTm`, `On Scene DtTm`

-- Query 1: Average response delay per Call Type & City (parses
timestamps inline)
SELECT
    `Call Type`,
    City,
    AVG(
        unix_timestamp(`On Scene DtTm`, 'M/d/yyyy h:mm:ss a') -
        unix_timestamp(`Response DtTm`, 'M/d/yyyy h:mm:ss a')) / 60.0
    ) AS avg_delay_min,
    COUNT(*) AS count_calls
FROM fire_calls
GROUP BY `Call Type`, City
ORDER BY avg_delay_min DESC
LIMIT 10;

-- Query 2: Count of "slow" responses (>10 min) per City (again
computed inline)
```

```
SELECT
    City,
    COUNT(*) AS slow_calls
FROM fire_calls
WHERE
    (unix_timestamp(`On Scene DtTm`, 'M/d/yyyy h:mm:ss a') -
     unix_timestamp(`Response DtTm`, 'M/d/yyyy h:mm:ss a')) / 60.0 >
    10
GROUP BY City
ORDER BY slow_calls DESC
LIMIT 10;
```

▶ _sqlfd: pyspark.sql.connect.dataframe.DataFrame = [City: string, slow_calls: long]

Table

ⓘ This result is stored as `_sqlfd` and can be used in other Python and SQL cells.

```
%sql
--  OPTIMIZED (works with the columns: uses Response DtTm, not Received DtTm)
```

```
-- Assumes a temp view `fire_calls` with: `Call Type`, City,
`Response DtTm`, `On Scene DtTm`

-- 0) Stage: precompute delay & apply EARLY filters (narrow
columns)
CREATE OR REPLACE TEMP VIEW fire_calls_stage AS
SELECT
    `Call Type` AS call_type,
    City      AS city,
    (
        unix_timestamp(`On Scene DtTm`, 'M/d/yyyy h:mm:ss a') -
        unix_timestamp(`Response DtTm`, 'M/d/yyyy h:mm:ss a')
    ) / 60.0 AS delay_min
FROM fire_calls
WHERE City = 'San Francisco'          -- early filter
    AND `Call Type` IS NOT NULL;       -- drop null group keys

CREATE OR REPLACE TEMP VIEW fire_calls_stage_part AS
SELECT /*+ REPARTITION(16, call_type) */ -- tune 16 to your
cluster size
    call_type, city, delay_min
FROM fire_calls_stage
WHERE delay_min > 0;                  -- prune before wide ops

-- Query A: Average delay per Call Type (aggregate after pruning)
SELECT
    call_type,
    COUNT(*)           AS count_calls,
    ROUND(AVG(delay_min),2) AS avg_delay_min
FROM fire_calls_stage_part
GROUP BY call_type
ORDER BY avg_delay_min DESC
LIMIT 10;

-- Query B: Count of "slow" responses (>10 min) per City
SELECT
    city,
    COUNT(*) AS slow_calls
FROM fire_calls_stage_part
WHERE delay_min > 10
GROUP BY city
```

```
ORDER BY slow_calls DESC  
LIMIT 10;
```

▶ _sqlfd: pyspark.sql.connect.dataframe.DataFrame = [city: string, slow_calls: long]

Table

This result is stored as `_sqlfd` and can be used in other Python and SQL cells.

```
from pyspark.sql.functions import col, avg, count, desc

# (City/Call Type filtered, response_delay_min computed, and > 0)

# 1) Keep only columns needed before wide ops (narrow -> fewer
bytes to shuffle)
df_narrow = df_clean.select("Call Type", "response_delay_min")

# 2) Partition by the group key to reduce shuffle during the
aggregation
df_partitioned = df_narrow.repartition(16, "Call Type") # adjust
16 to your cluster size

# 3) Single wide op: groupBy + agg
df_agg = (
    df_partitioned
    .groupBy("Call Type")
    .agg(
```

```

        avg("response_delay_min").alias("avg_delay_min"),
        count("*").alias("count_calls")
    )
)

# 4) Global sort only after aggregation (on many fewer rows)
top10 = df_agg.orderBy(desc("avg_delay_min")).limit(10)

print("✅ PySpark – Optimized (filters early + partition by key +
minimal shuffles, no persist):")
top10.show(truncate=False)

```

- ▶ 📄 df_agg: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_delay_min: double ... 1 more field]
 - ▶ 📄 df_narrow: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, response_delay_min: double]
 - ▶ 📄 df_partitioned: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, response_delay_min: double]
 - ▶ 📄 top10: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_delay_min: double ... 1 more field]
- ✅ PySpark – Optimized (filters early + partition by key + minimal shuffle s, no persist):

Call Type	avg_delay_min	count_calls
Mutual Aid / Assist Outside Agency	86.18238095238097	21
High Angle Rescue	11.198544061302687	261
Watercraft in Distress	11.052735849056605	106
Water Rescue	8.801310844464568	4417
Train / Rail Incident	8.619045801526717	262
Suspicious Package	7.456808510638298	47
Marine Fire	7.300263157894739	38
Medical Incident	6.5862889263119415	921003
Other	5.812180554645021	15253
Train / Rail Fire	5.754761904761905	21

```
# --- 1) Discover catalogs/schemas available in this workspace ---
```

```
print("Available catalogs:")
spark.sql("SHOW CATALOGS").show(truncate=False)

current_catalog = spark.sql("SELECT current_catalog()").first()[0]
current_schema = spark.sql("SELECT current_schema()").first()[0]
print(f"current_catalog={current_catalog}, current_schema={current_schema}")
```

Available catalogs:

```
+-----+
|catalog |
+-----+
|samples |
|system  |
|workspace|
+-----+
```

current_catalog=workspace, current_schema=default

```
#  Use the workspace catalog and default schema (supported on
# serverless)
spark.sql("USE CATALOG workspace")
spark.sql("USE SCHEMA default")

# Create a Volume once (safe to re-run)
spark.sql("CREATE VOLUME IF NOT EXISTS analytics_vol COMMENT
'Analysis outputs'")

# Write your results to the Volume
dest_path =
"/Volumes/workspace/default/analytics_vol/fire_calls_top10_parquet
"

# Save as Parquet
top10.write.mode("overwrite").parquet(dest_path)

print("✅ Successfully wrote results to:", dest_path)
```

✓ Successfully wrote results to: /Volumes/workspace/default/analytics_vo
l/fire_calls_top10_parquet

```
# Actions vs Transformations
from pyspark.sql.functions import col, to_timestamp, round
import time

TS_FMT = "M/d/yyyy h:mm:ss a"
df = df_spark_select

# ---- Transformations (lazy) ----
t0 = time.time()
t = (
    df
    .filter(col("City") == "San Francisco")
    .filter(col("Call Type").isNotNull())
    .withColumn("received_ts", to_timestamp(col("Response DtTm"),
TS_FMT))
    .withColumn("on_scene_ts", to_timestamp(col("On Scene DtTm"),
TS_FMT))
    .withColumn(
        "response_delay_min",
        round((col("on_scene_ts").cast("long") -
col("received_ts").cast("long")) / 60.0, 2)
    )
    .select("Call Type", "City", "response_delay_min")
)
print(f"Transformations built in {time.time() - t0:.4f}s (no Spark
job yet)")

# ---- Actions (eager) ----
print("\nAction 1: count()")
t1 = time.time()
cnt = t.count()      # triggers a job
print(f"count() = {cnt:,} rows, took {time.time() - t1:.2f}s")

print("\nAction 2: show()")
t2 = time.time()
t.show(5, truncate=False) # triggers another job
```

```
print(f"show() took {time.time() - t2:.2f}s")

print("\n🔍 Each action re-executes the transformations unless you
materialize it.")

# localCheckpoint() breaks lineage and materializes the result in
executor memory/disk
# without using metastore features that are blocked on serverless.
t_ckpt = t.localCheckpoint(eager=True)

print("\nAction 3 (after local checkpoint): show()")
t3 = time.time()
t_ckpt.show(5, truncate=False) # should avoid recomputing full
upstream lineage
print(f"show() after local checkpoint took {time.time() -
t3:.2f}s")
```

▶ 📄 df: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 2 more fields]

▶ 📄 t: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 1 more field]

▶ 📄 t_ckpt: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 1 more field]

```
|Alarms      |San Francisco|NULL          |
|Alarms      |San Francisco|0.0           |
|Structure Fire|San Francisco|2.4          |
+-----+-----+-----+
```

only showing top 5 rows

show() took 0.27s

🔍 Each action re-executes the transformations unless you materialize it.

Action 3 (after local checkpoint): show()

```
+-----+-----+-----+
|Call Type |City       |response_delay_min|
+-----+-----+-----+
|Alarms    |San Francisco|2.7          |
|Alarms    |San Francisco|NULL         |
|Alarms    |San Francisco|NULL         |
|Alarms    |San Francisco|0.0          |
|Structure Fire|San Francisco|2.4          |
+-----+-----+-----+
```

```
+-----+-----+
only showing top 5 rows
```

```
# MLLib – Binary Classification on Fire Calls (serverless-safe)
from pyspark.sql.functions import col, when, to_timestamp, hour,
dayofweek
from pyspark.ml import Pipeline
from pyspark.ml.feature import StringIndexer, OneHotEncoder,
VectorAssembler
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.evaluation import BinaryClassificationEvaluator,
MulticlassClassificationEvaluator

TS_FMT = "M/d/yyyy h:mm:ss a"

# Use your selected columns DataFrame from earlier
df_src = df_spark_select

# Pick the available "start" timestamp column
start_ts_col = "Received DtTm" if "Received DtTm" in
df_src.columns else "Response DtTm"

# 1) Prepare label & features
df_ml = (
    df_src
    .withColumn("start_ts", to_timestamp(col(start_ts_col),
TS_FMT))
    .withColumn("on_scene_ts", to_timestamp(col("On Scene DtTm")),
TS_FMT))
    .withColumn("response_delay_min",
(col("on_scene_ts").cast("long") -
col("start_ts").cast("long"))/60.0)
    .withColumn("slow_response", when(col("response_delay_min") >=
10, 1).otherwise(0))
    .withColumn("hr", hour(col("start_ts")))
    .withColumn("dow", dayofweek(col("start_ts")))
    .dropna(subset=["slow_response", "hr", "dow", "City", "Call
Type"])
)
```

```
# Categorical -> index -> one-hot
cat_cols = ["Call Type", "City"]
idx_cols = [c + "_idx" for c in cat_cols]
ohe_cols = [c + "_vec" for c in cat_cols]

indexers = [StringIndexer(inputCol=c, outputCol=i,
handleInvalid="keep") for c, i in zip(cat_cols, idx_cols)]
encoder = OneHotEncoder(inputCols=idx_cols, outputCols=ohe_cols)
assembler = VectorAssembler(inputCols=ohe_cols + ["hr", "dow"],
outputCol="features")

# 2) Model
lr = LogisticRegression(labelCol="slow_response",
featuresCol="features", maxIter=50)

pipeline = Pipeline(stages=indexers + [encoder, assembler, lr])

# 3) Train / Test
train_df, test_df = df_ml.randomSplit([0.8, 0.2], seed=42)
model = pipeline.fit(train_df)
pred = model.transform(test_df)

# 4) Evaluate
auc = BinaryClassificationEvaluator(labelCol="slow_response",
rawPredictionCol="rawPrediction").evaluate(pred)
f1 = MulticlassClassificationEvaluator(labelCol="slow_response",
predictionCol="prediction", metricName="f1").evaluate(pred)
print(f"✓ Test AUC: {auc:.4f}")
print(f"✓ Test F1 : {f1:.4f}")

# Inspect
pred.select("Call
Type", "City", "response_delay_min", "slow_response", "prediction", "pr
obability").show(10, truncate=False)
print("\nConfusion matrix (prediction vs label):")
(pred.groupBy("slow_response", "prediction").count().orderBy("slow_
response", "prediction")).show()
```

- ▶ df_ml: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 8 more fields]
- ▶ df_src: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 2

more fields]

- ▶ pred: pyspark.sql.connect.dataframe.DataFrame
- ▶ test_df: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 8 more fields]
- ▶ train_df: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 8 more fields]

Aircraft Emergency SFO	NULL	0	0.0	
[0.5992085342183672, 0.4007914657816328]				
Aircraft Emergency SFO	12.6	1	0.0	
[0.5875709822564589, 0.4124290177435411]				
Alarms DC	NULL	0	0.0	
[0.98835964573658, 0.011640354263420027]				
Alarms Fort Mason	5.75	0	0.0	
[0.9806818508422595, 0.01931814915774055]				
Alarms Fort Mason	4.633333333333334	0	0.0	
[0.9807713700576458, 0.019228629942354192]				
Alarms Fort Mason	6.35	0	0.0	
[0.9807713700576458, 0.019228629942354192]				
Alarms Fort Mason	3.9166666666666665	0	0.0	
[0.9807713700576458, 0.019228629942354192]				
Alarms Fort Mason	4.883333333333334	0	0.0	
[0.980690821117373, 0.019309178882626954]				
Alarms Fort Mason	5.2666666666666667	0	0.0	
[0.980037387363936, 0.019962612636064025]				
Alarms Fort Mason	4.35	0	0.0	
[0.9805633327841002, 0.01943666721589976]				

top10.explain()

```
S LAST], partitionOrderCount=0)
    +- PhotonGroupingAgg(keys=[Call Type#15921], functions=[avg(response_delay_min#19359), count(1)])
        +- PhotonShuffleExchangeSource
            +- PhotonShuffleMapStage REPARTITION_BY_NUM,
[id=#16701]
            +- PhotonShuffleExchangeSink hashpartitioning(Call Type#15921, 16)
                +- PhotonProject [Call Type#15921, roun
```

```
d((cast((cast(gettimestamp(On Scene DtTm#15928, M/d/yyyy h:mm:ss a, Time
stampType, try_to_timestamp, Some(Etc/UTC), true) as bigint) - cast(gett
imestamp(Response DtTm#15927, M/d/yyyy h:mm:ss a, TimestampType, try_to_
timestamp, Some(Etc/UTC), true) as bigint)) as double) / 60.0), 2) AS re
sponse_delay_min#19359]
+- PhotonFilter (((isnotnull(City#15
934) AND isnotnull(Call Type#15921)) AND (City#15934 = San Francisco)) A
ND (round((cast((cast(gettimestamp(On Scene DtTm#15928, M/d/yyyy h:mm:ss
a, TimestampType, try_to_timestamp, Some(Etc/UTC), true) as bigint) - ca
st(gettimestamp(Response DtTm#15927, M/d/yyyy h:mm:ss a, TimestampType,
try_to_timestamp, Some(Etc/UTC), true) as bigint)) as double) / 60.0),
2) > 0.0))
```

```
top10.limit(10).show()
```

Call Type	avg_delay_min	count_calls
Mutual Aid / Assi...	86.18238095238097	21
High Angle Rescue	11.198544061302687	261
Watercraft in Dis...	11.052735849056605	106
Water Rescue	8.801310844464568	4417
Train / Rail Inci...	8.619045801526717	262
Suspicious Package	7.456808510638298	47
Marine Fire	7.300263157894739	38
Medical Incident	6.5862889263119415	921003
Other	5.812180554645021	15253
Train / Rail Fire	5.754761904761905	21

```
from pyspark.sql.functions import col
import time

df_src = top10
```

```
start = time.time()
_ = df_src.count() # first action (full recompute)
t1 = time.time() - start

start = time.time()
_ = df_src.count() # second action (recomputes
again)
t2 = time.time() - start

print(f"Baseline (no materialization) - first count: {t1:.2f}s")
print(f"Baseline (no materialization) - second count: {t2:.2f}s")

# localCheckpoint()
start = time.time()
df_ckpt = df_src.localCheckpoint(eager=True) # materialize now
t3 = time.time() - start
print(f"\nMaterialize with localCheckpoint(eager=True): {t3:.2f}s"

# First action after checkpoint
start = time.time()
_ = df_ckpt.count()
t4 = time.time() - start

# Second action after checkpoint
start = time.time()
_ = df_ckpt.count()
t5 = time.time() - start

print(f"After checkpoint - first count: {t4:.2f}s")
print(f"After checkpoint - second count: {t5:.2f}s")

start = time.time()
df_ckpt.show(5, truncate=False)
print(f"show() from checkpoint took: {time.time() - start:.2f}s")
```

- ▶ df_ckpt: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_delay_min: double ... 1 more field]
- ▶ df_src: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_delay_min: double ... 1 more field]

```
Baseline (no materialization) - first count: 4.74s
Baseline (no materialization) - second count: 4.12s
```

```
Materialize with localCheckpoint(eager=True): 4.99s
```

```
After checkpoint - first count: 0.24s
```

```
After checkpoint - second count: 0.12s
```

Call Type	avg_delay_min	count_calls
Mutual Aid / Assist Outside Agency	86.18238095238097	21
High Angle Rescue	11.198544061302687	261
Watercraft in Distress	11.052735849056605	106
Water Rescue	8.801310844464568	4417
Train / Rail Incident	8.619045801526717	262

only showing top 5 rows

show() from checkpoint took: 0.22s

- ▶ check_df: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_delay_min: double ... 1 more field]
- ▶ df_clean: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 5 more fields]
- ▶ df_filtered: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 2 more fields]
- ▶ df_result: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, avg_delay_min: double ... 1 more field]
- ▶ df_with_delay: pyspark.sql.connect.dataframe.DataFrame = [Call Type: string, City: string ... 5 more fields]

Medical Incident	6.586288926311867	921003
Other	5.812180554644988	15253
Train / Rail Fire	5.754761904761905	21

only showing top 10 rows

Successfully wrote results to: /Volumes/workspace/default/analytics_v ol/fire_calls_top10_parquet

Output successfully reloaded:

Call Type	avg_delay_min	count_calls
Mutual Aid / Assist Outside Agency	86.18238095238097	21